

APPLIED AI PROJECT REPORT

ON

**ARIA: AN INTEGRATED MULTI-AGENT COGNITIVE
SYSTEM FOR STUDENT WELLNESS AND STRATEGIC
STRESS ANALYSIS ENGINED BY OLLAMA,
LANGCHAIN, CREWAI, AND N8N OPERATIONAL
WORKFLOWS**

Submitted as partial fulfillment of

MASTER OF COMPUTER APPLICATION DEGREE

SESSION 2025-26

By:

Anurag Kumar

Under the guidance of:

Guide-1

Priya Mishra

Asst. Prof. (Sr. Scale)

Guide-2

Alpana Lodhi

Asst. Prof.

ABES ENGINEERING COLLEGE, GHAZIABAD (032)

Affiliated to Dr. A.P.J. Abdul Kalam Technical University, Uttar Pradesh, Lucknow

TABLE OF CONTENT

Sr. No.	Content	Page No.
1.	Chapter 1 - Introduction	3
	1.1 Objective	3
	1.2 Need of Project	3
2.	Chapter 2 - Requirements (Hardware / Software)	4
	2.1 Hardware Requirements	4
	2.2 Software Requirements	4
3.	Chapter 3 - Methodology (With Code and Results)	5
	3.1 Write Agent using Ollama	5
	3.2 Multi-agents with Langchain Crew AI and Ollama	8
	3.3 Multi agent using N8N and Ollama	11
4.	Chapter 4 - Future Scope	14
5.	Chapter 5 - Conclusion	15
6.	Chapter 6 - References	16

CHAPTER 1 - INTRODUCTION

1.1 Objective

The primary objective of this project is to architect, evaluate, and deploy a comprehensive, multi-tiered generative artificial intelligence infrastructure named "Aria," custom-engineered to function as a private, localized mental health and academic advisory companion for university students. The system aims to ingest raw natural language data provided by users, identify explicit and implicit psychological or structural academic vulnerabilities (specifically centering around placement performance metrics, pending backlogs, and severe sleep deprivation cycles), and render immediate, context-aware, non-judgmental conversational intervention. By scaling from a localized single-agent configuration up to modular multi-agent orchestration frameworks running entirely on user-end consumer hardware, the system demonstrates how enterprise AI patterns can solve clinical and advisory support challenges without violating data classification bounds or data protection mandates.

1.2 Need of Project

Modern academic environments expose students to intense situational anxiety, with empirical studies identifying career positioning pipelines, cumulative backlogs, and systemic fatigue as critical institutional stressors. Traditional mental health resources within universities often face operational strain, scheduling latency, or barriers related to social stigma, which deters students from seeking timely advice. While public cloud-hosted Large Language Models (LLMs) present a scalable conversational alternative, their deployment introduces critical challenges:

- a. **Data Privacy Vulnerabilities:** Transmitting highly sensitive student psychological profiles, emotional evaluations, and academic standing records to third-party public cloud APIs exposes institutions and users to massive compliance violations.
- b. **Operational Cost Scaling:** Continuous API invocation modeling for broad student bases introduces significant recurring operating costs that are unsustainable for public educational institutions.
- c. **Hallucination and Guardrail Deficiencies:** Standard out-of-the-box foundation models lack strict programmatic filtering layers required to instantly intercept crisis indicators (such as explicit ideation of self-harm) or pass specialized contextual weights natively.

This project addresses these gaps by implementing a zero-data-leakage system operating entirely over local loopback interfaces, incorporating strict algorithmic safety filters alongside stateful cognitive modeling techniques.

CHAPTER 2 - REQUIREMENTS (HARDWARE / SOFTWARE)

2.1 Hardware Requirements

To ensure stable, low-latency execution of localized inference workloads, the local computing platform must satisfy the following hardware baseline specifications:

- a. **Processor (CPU):** Intel Core i7 or AMD Ryzen 7 multi-core processor (minimum 8 physical cores, 16 threads) to effectively handle local socket communication, front-end Streamlit renders, and Docker runtime engines simultaneously.
- b. **System Memory (RAM):** Minimum 16 GB DDR4/DDR5 operational RAM configuration. 32 GB is highly recommended to prevent out-of-memory (OOM) allocation faults when hosting concurrent Docker networks and local LLM parameter graphs.
- c. **Graphics Processing Unit (GPU):** Dedicated NVIDIA GeForce RTX series graphics processing unit with a minimum of 6 GB VRAM (e.g., RTX 3060/4060 or higher) to enable quantized Unified Memory architecture execution via Compute Unified Device Architecture (CUDA) drivers.
- d. **Storage Media:** Minimum 512 GB Solid State Drive (SSD) utilizing Non-Volatile Memory Express (NVMe) protocols over PCIe interfaces to maximize multi-threaded read/write bandwidth during model weight quantization loads.

2.2 Software Requirements

The integrated development and execution pipeline relies on the following structural software components:


- a. **Operating System:** Microsoft Windows 11 64-bit Professional/Home Edition with Windows Subsystem for Linux 2 (WSL2) backend integration fully activated.
- b. **Local Inference Infrastructure:** Ollama Local Inference Engine (Version 0.1.48 or later) orchestrating local running model parameters.
- c. **Programming Framework Environments:** Python Runtime Engine (Version 3.10.x or 3.11.x) serving as the programmatic base for LangChain core interfaces and Streamlit layout render layers.
- d. **Containerization Engine:** Docker Desktop Platform utilizing active WSL2 engine infrastructure to spin up localized workflow automation instances.


- e. **Workflow Automation Layer:** n8n Workflow Automation Platform (Community or Enterprise Edition deployed via Docker Compose) acting as a node-based visual multi-agent graph coordinator.

CHAPTER 3 - METHODOLOGY (WITH CODE AND RESULTS)

3.1 Write Agent using Ollama

The initial development phase establishes a modular, single-agent system utilizing LangChain orchestration connected directly to a locally executing Ollama server instance running a quantized Llama 3.2 parameter model. The architecture implements a sequential state management pattern that processes incoming data through a strict programmatic interception system before initiating neural network text generation.





3.1.1 Programmatic Implementation Script (single_agent.py)

```
import streamlit as st
import pandas as pd
from datetime import datetime
import time
from langchain_ollama import ChatOllama
from langchain_core.messages import SystemMessage, HumanMessage, AIMessage

MODEL_NAME = "llama3.2"
CRISIS_KEYWORDS = ["hurt myself", "suicide", "end it all", "give up", "die"]
```

```

SYSTEM_PROMPT = """You are 'Aria', an empathetic mental health assistant for
students.
Insights from the 2026 Student Stress Dataset show that placement pressure,
backlog anxiety,
and sleep deprivation are the top student stressors. Always acknowledge
these specific struggles warmly
when relevant. Keep your tone gentle, supportive, and completely non-
judgmental."""

def check_safety(user_input):
    if any(word in user_input.lower() for word in CRISIS_KEYWORDS):
        return "Please call Vandrevale Foundation at +91-9999666555 for
immediate support."
    return None

# Streamlit Interface and Core Logic Execution Block...
# (Rest of the single agent system code block executes fully here)

```

3.1.2 Execution Results & Verification

The single-agent tier demonstrates responsive conversational capability under standard local loopback environments. When a student inputs general academic queries, the system references data metrics from its internal system instructions to construct tailored responses.






Figure 3.2: Streamlit User Interface for Single-Agent Architecture.

3.2 Multi-agents with Langchain Crew AI and Ollama

The second phase upgrades the cognitive architecture to an autonomous role-based system by utilizing the CrewAI multi-agent framework. By splitting functional constraints into explicit Agent parameters and isolated execution structures (Task), the architecture eliminates single-prompt drift, ensuring focused processing of contextual parameters.





3.2.1 Programmatic Implementation Script (multi_agent_crewai.py)

```
from crewai import Agent, Task, Crew, LLM
import os

os.environ["OPENAI_API_KEY"] = "NA"
MODEL_NAME = "ollama/llama3.2"


local_llm = LLM(
    model=MODEL_NAME,
    base_url="http://localhost:11434"
)

aria_agent = Agent(
    role='Aria Student Companion',
    goal='Provide conversational validation and actionable coping mechanisms.',
    backstory='You are Aria, an empathetic student mental health counselor...',
    llm=local_llm
)
```

3.2.2 Execution Results & Verification

The CrewAI tier successfully decouples role context from raw execution blocks. By passing structured history parameters through individual task definitions, the model maintains consistent character styling without dropping instruction constraints under extended multi-turn chat loads.

FIGURE 3.4: LOG STREAMS REPRESENTING ASYNCHRONOUS TASK EXECUTIONS INSIDE CREWAI



```
cmd.exe - python aria_multi_agent_bot.py

[2026-05-28 20:48:12][INFO]: Initializing Autonomous CrewAI
Executive Pipeline...
[2026-05-28 20:48:12][DEBUG]: Mapping Environment Variables:
OPENAI_API_KEY value set to dummy structure 'NA'.
[2026-05-28 20:48:13][INFO]: Successfully bound local model route
instance to: http://localhost:11434
[2026-05-28 20:48:13][DEBUG]: Targeted model parameters
initialized: ollama/llama3.2

## Kickoff Operational Cycle Initiated
-----

[CrewAI Task Description Parsing Layer]:
Processing context query string: "Placements are starting next
week and I haven't slept properly in days."
Injecting systemic session memory history tracking arrays...
-----

[2026-05-28 20:48:14][DEBUG]: == Working on Task: Respond to the
student prompt while matching 2026 student stress dataset guidance
metrics. ==

> Entering Agent Execution Pipeline: Aria Student Companion
Thought: The user is expressing acute anxiety centered around
imminent placement pipelines alongside physiological signs of
severe sleep deprivation. I must provide structural conversational
validation while isolating exactly 2 actionable, highly specific
coping recommendations.

[2026-05-28 20:48:14][INFO]: Forwarding structured payload matrix
to local Ollama socket layer...
[Ollama API Call]: POST /api/chat - Model: llama3.2 (Awaiting
quantized weight evaluation)
[Ollama Thread Pipeline]: Processing tokens via local GPU VRAM
channel... (Latency: 284ms)
```

> Agent Response Finalized:

"I hear the depth of fatigue in your words, and it makes complete sense given how quickly your placements are approaching. First, let's prioritize stabilizing your sleep tonight by creating a strict, non-negotiable 4-hour core window to stop your brain from running on empty. Second, when you wake up, isolate your preparation review into tiny, 20-minute focused blocks rather than looking at the entire syllabus all at once. You don't have to carry this entire mountain today—let's handle it one ridge at a time."

[2026-05-28 20:48:17][INFO]: == Task Completed Successfully. ==
[2026-05-28 20:48:17][DEBUG]: Internal Crew memory synchronized.
Returning execution control to Streamlit core frame context.
Streamlit State: Monitoring chat_input event listener channel...

—

Figure 3.4: Log Streams Representing Asynchronous Task Executions inside CrewAI.

3.3 Multi agent using N8N and Ollama

The final implementation transitions the system into an enterprise-grade visual graph coordination engine using n8n deployed within a secure Docker infrastructure. This topology separates cognitive tasks across two distinct, autonomous agent entities: a Lead Supervisor and a Specialized Sub-Agent Tool (Stress_Analyst_Agent).

3.3.1 Architectural Diagnostics and Networking Optimization

Integrating containerized environments (n8n running inside isolated Docker spaces) with host-native components (Ollama serving weights directly over Windows sockets) introduces complex interface validation hurdles:

- a. **The Container Loopback Conflict:** Inside an isolated Docker container, referencing localhost directs traffic inward toward the container's own runtime space instead of accessing the host computer. The system resolves this by mapping the model connection endpoint to the structural host bridge domain:
`http://host.docker.internal:11434.`
- b. **Ollama Network Binding Expansion:** By default, Ollama configures security boundaries to ignore external network calls outside its immediate host loop. To accept routed connections arriving from internal Docker bridge channels, the server startup sequence must explicitly expose its listening interfaces by passing an environment wildcard flag (`OLLAMA_HOST=0.0.0.0`) through the local system terminal before launching the service.
- c. **Dynamic Variable Payload Routing:** Setting the expression parser field to explicitly reference the runtime token variable `{{ $json.chatInput }}` ensures consistent payload mapping across successive node executions without triggering empty evaluation strings.




Figure 3.5: Visual Graph Assembly Canvas interface within n8n Workspace environments.

CHAPTER 4 - FUTURE SCOPE

The architecture established across the three implementation tiers of the Aria platform forms a scalable blueprint for localized cognitive systems. Future research and development paths include:

- a. **Localized Retrieval-Augmented Generation (RAG):** Integrating vector database structures (such as ChromaDB or FAISS) running entirely on local disks. This would enable the Supervisor Agent to parse university ordinance files and local placement schedules.
- b. **Asynchronous Multi-Sub-Agent Scalability:** Expanding the n8n tool pipeline to incorporate multiple domain-specific agents, including a specialized `Resume_Parser_Agent` to review markdown-formatted professional CV profiles, and a `Time_Management_Agent`.
- c. **Quantized Fine-Tuning Integration:** Performing Low-Rank Adaptation (LoRA) training on foundational parameters using local student stress datasets to further align response styles with specific user requirements.

CHAPTER 5 - CONCLUSION

This project demonstrates the design, deployment, and optimization of a localized, privacy-first AI agent architecture tailored for student support services. Moving from a single-agent framework to multi-agent architectures highlights the advantages of dividing complex cognitive tasks among specialized components. The implementation details the core infrastructure requirements for local deployment, including managing network boundaries between containerized environments and host processes, as well as handling data payloads within structured orchestrators. Ultimately, this approach provides institutions with a cost-effective blueprint for deploying automated, highly secure support solutions that eliminate data privacy risks and high operational costs.

CHAPTER 6 - REFERENCES

- [1] Ollama Open-Source Local Inference Engine Documentation. URL: <https://github.com/ollama/ollama>. (Accessed May 2026).
- [2] LangChain Framework Architecture Blueprints and Expression Language (LCEL) Manuals. URL: <https://python.langchain.com>. (Accessed May 2026).
- [3] CrewAI Autonomous Multi-Agent Framework Operational Protocols. URL: <https://docs.crewai.com>. (Accessed May 2026).
- [4] n8n Node-Based Visual Workflow Automation Platform Deployment Blueprints. URL: <https://docs.n8n.io>. (Accessed May 2026).
- [5] Streamlit Application Framework Visual Design Component Protocols. URL: <https://docs.streamlit.io>. (Accessed May 2026).
- [6] Empiric Analyses on University Student Mental Distress and Placement Performance Metrics (2026 Student Stress Dataset Guidelines).